

A4

Kerui Du

05/10/2021

Q1

a

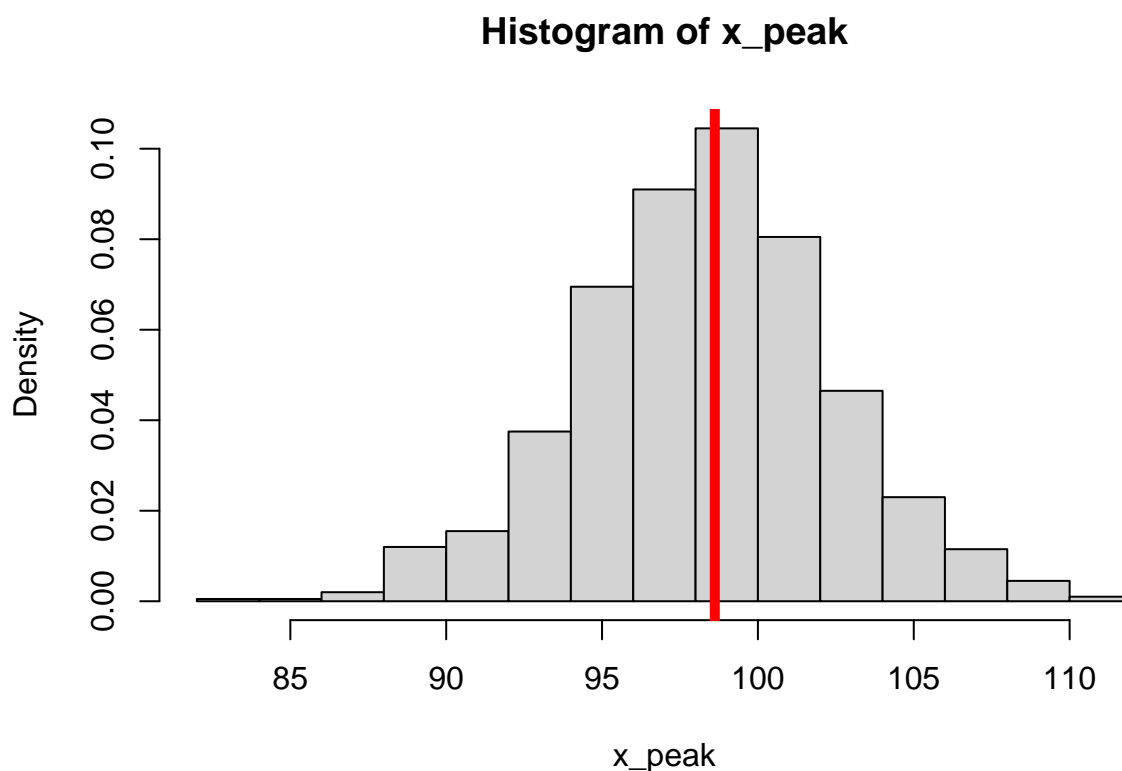
```
Caffeine2.df=data.frame(n=rep(300,5),levels=seq(0,200,50),
                        A_grades=c(109,155,175,158,103))
fit=glm(cbind(A_grades,n-A_grades)~levels+I(levels^2),family=binomial,data = Caffeine2.df)

ns=Caffeine2.df$n
xs=Caffeine2.df$levels
prob = with(Caffeine2.df,A_grades/n)

est = matrix(0,nr=1000,nc=2)
x_peak = -coef(fit)[2]/coef(fit)[3]/2

for (i in 1:nrow(est)) {
  ys = rbinom(length(ns),ns,prob)
  model=glm(cbind(ys,ns-ys)~xs+I(xs^2),family=binomial)
  est[i,]=coef(model)[2:3]
}

hist(-est[,1]/est[,2]/2,probability = T,xlab = 'x_peak',main='Histogram of x_peak')
abline(v=x_peak,lwd=5,col='red')
```



x\_peak looks normality distributed and the peak of the histogram is around 98, which is pretty close to the x\_peak(red line).

1b

```
c(2*x_peak-quantile(-est[,1]/est[,2]/2,.975),
  2*x_peak-quantile(-est[,1]/est[,2]/2,.025))
```

```
##      levels      levels
## 90.61776 107.43167
```

1c

The interval we obtained above is similar to the result that we gain from assignment 3, and it seems that the non-parametric CIs are slightly wider than the parametric CIs.

Q2

2a

```
train.df = read.table("data/train.txt")
names(train.df) = c("D", paste("V", 1:256, sep=""))
all(!is.na(train.df)) # examine if NA exist
```

```
## [1] TRUE
```

```
all(train.df$D==7|train.df$D==3) # D has value either 7 or 3
```


```
## [1] TRUE
```

```
all(train.df[-1]<=1 & train.df[-1]>=-1) # rest of columns between -1 to 1
```

```
## [1] TRUE
```

## 2b

```
mod=par(mfrow=c(5,5), mar = c(1,1,1,1))
for(k in 1:25){
  z = matrix(unlist(train.df[k,-1]), 16,16)
  zz = z
  for(j in 16:1)zz[,j]=z[,17-j]
  image(zz, col = gray((32:0)/32))
  box()
  text(0.1,0.9,train.df$D[k],cex=1.5)
}
```



```
par(mod)
```

By looking these image, we can easily detect that the there is a significant difference between digit 3 and digit 7 at bottom of each image, therefore the position v179/v180; v195/v196 should have value close to 1.

2c

```
corr = cor(train.df[,1:257])[,1]
varb=names(sort(abs(corr),decreasing = T))[2:21];varb
```

```
## [1] "V185" "V170" "V105" "V220" "V235" "V201" "V229" "V120" "V219" "V230"
## [11] "V104" "V189" "V205" "V169" "V234" "V121" "V204" "V186" "V173" "V221"
```

2d

```
Y = train.df$Y = ifelse(train.df$D==7,1,0)
train.df1 = data.frame(train.df[varb],Y)
Full.mod=glm(Y~.,family = binomial,data = train.df1)
fitted_logit = predict(Full.mod,train.df1)
head(fitted_logit)
```

```
##           1           2           3           4           5           6
##  9.250851 -14.077921 -10.735400  2.836090  5.742092  7.653596
```

2e

```
tb = table(round(train.df1$Y), round(fitted(Full.mod)));tb
```

```
##
##      0    1
##  0 651    7
##  1    7 638
```

```
sum(tb[tb!=diag(tb)])/sum(tb)
```

```
## [1] 0.01074444
```

2f

```
null.model = glm(Y~1, data=train.df, family=binomial)
#step(Full.mod, direction = "backward")
sub.mod = step(null.model, formula(Full.mod),direction="both", trace=0)
tb = table(round(train.df1$Y), round(fitted(sub.mod)));tb
```

```
##
##      0    1
##  0 650    8
##  1    9 636
```

```
sum(tb[tb!=diag(tb)])/sum(tb)
```

```
## [1] 0.01304682
```

The value of this model is pretty much same to the value from full.mod above.

2g

```
train.df2 = data.frame(train.df1[, names(sub.mod$coefficients)[-1]],Y)

set.seed(330)
PE.fun <- function() {
  sub <- sample(nrow(train.df2), floor(nrow(train.df2) * 0.9)) # 9 portions
  training <- train.df2[sub, ] # 90% for training
  testing <- train.df2[-sub, ] # 10% for testing
  glm.fit = glm(Y~.,data = training,family = binomial)
  tb=table(round(testing$Y),round(predict(glm.fit,testing[, -11],type = 'resp')))
  (tb[1,2]+tb[2,1])/sum(tb) # prediction error
}
PE.fun()
```

```
## [1] 0.01526718
```

The prediction error in cross validation depends on the observations that have been assigned in each training and testing dataset, then slightly different value from above, but it's close enough.

## 2h

```
test.df=read.table('data/test.txt')
Y=test.df$Y=ifelse(test.df$V1==7,1,0)

test1.df=data.frame(test.df[names(coef(sub.mod)[-1])],Y)
pr_out=table(test1.df$Y,round(predict(sub.mod,test1.df[, -11],type = 'resp')));pr_out
```

```
##
##      0    1
##  0 162    4
##  1   17  130
```

```
sum(pr_out[pr_out!=diag(pr_out)])/sum(pr_out)
```

```
## [1] 0.06709265
```

An estimate of the “out-of-sample” prediction error for the models from previous question is around 0.067, not much difference by ‘in-sample’ prediction error, which indicates both models are good for prediction.